

CS60075
Natural Language Processing
Autumn 2020

Module 6A

Sequence processing with Recurrent Networks

Sequences in NLP

Language is inherently temporal.

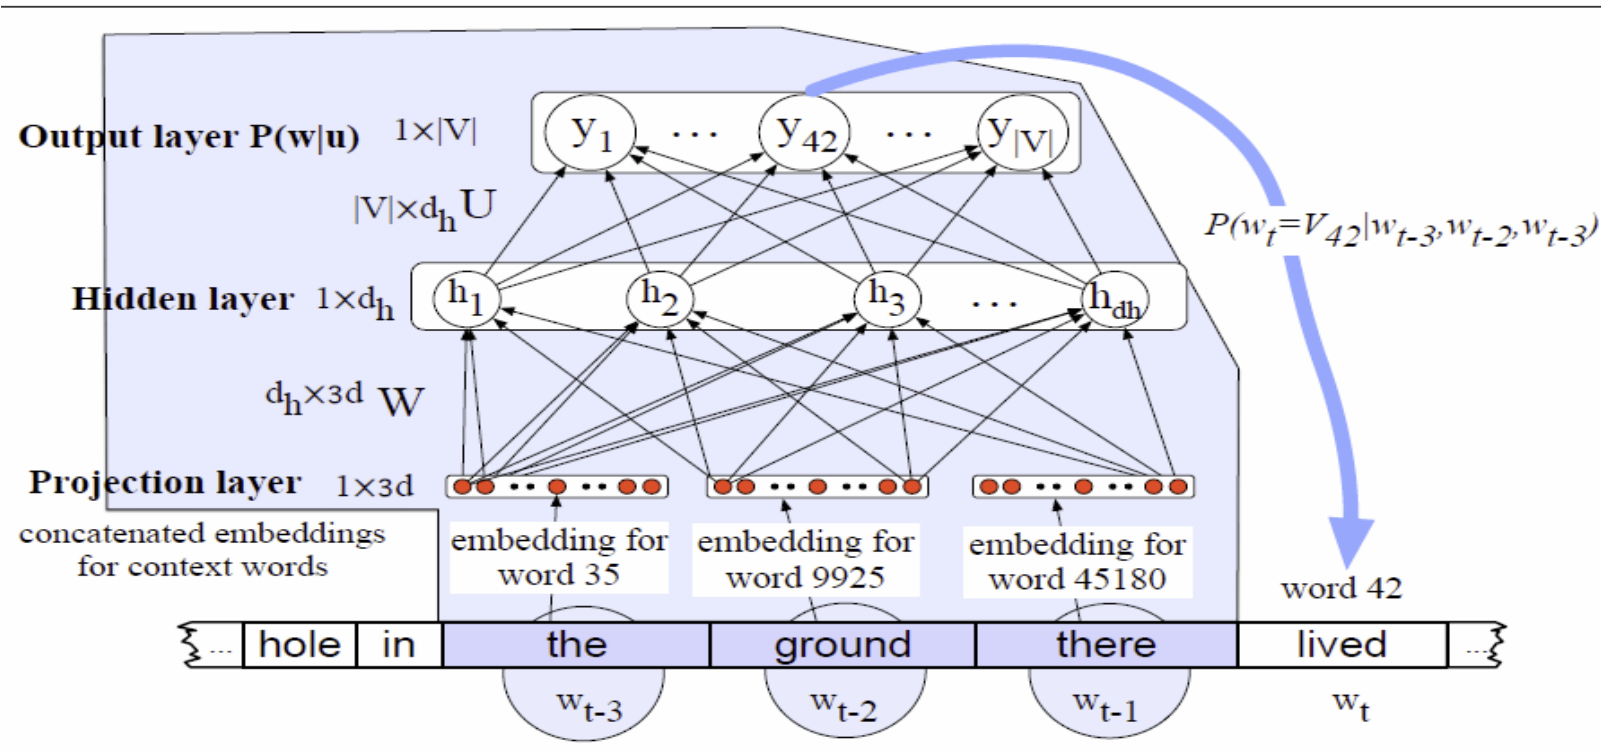
- I hope there are no more classes
- no I hope there are more classes

Applications of sequence processing

- Syntactic parsing
- Part of speech tagging
- Viterbi algorithm

3-gram Neural Language Model

Assume Pre-trained embeddings




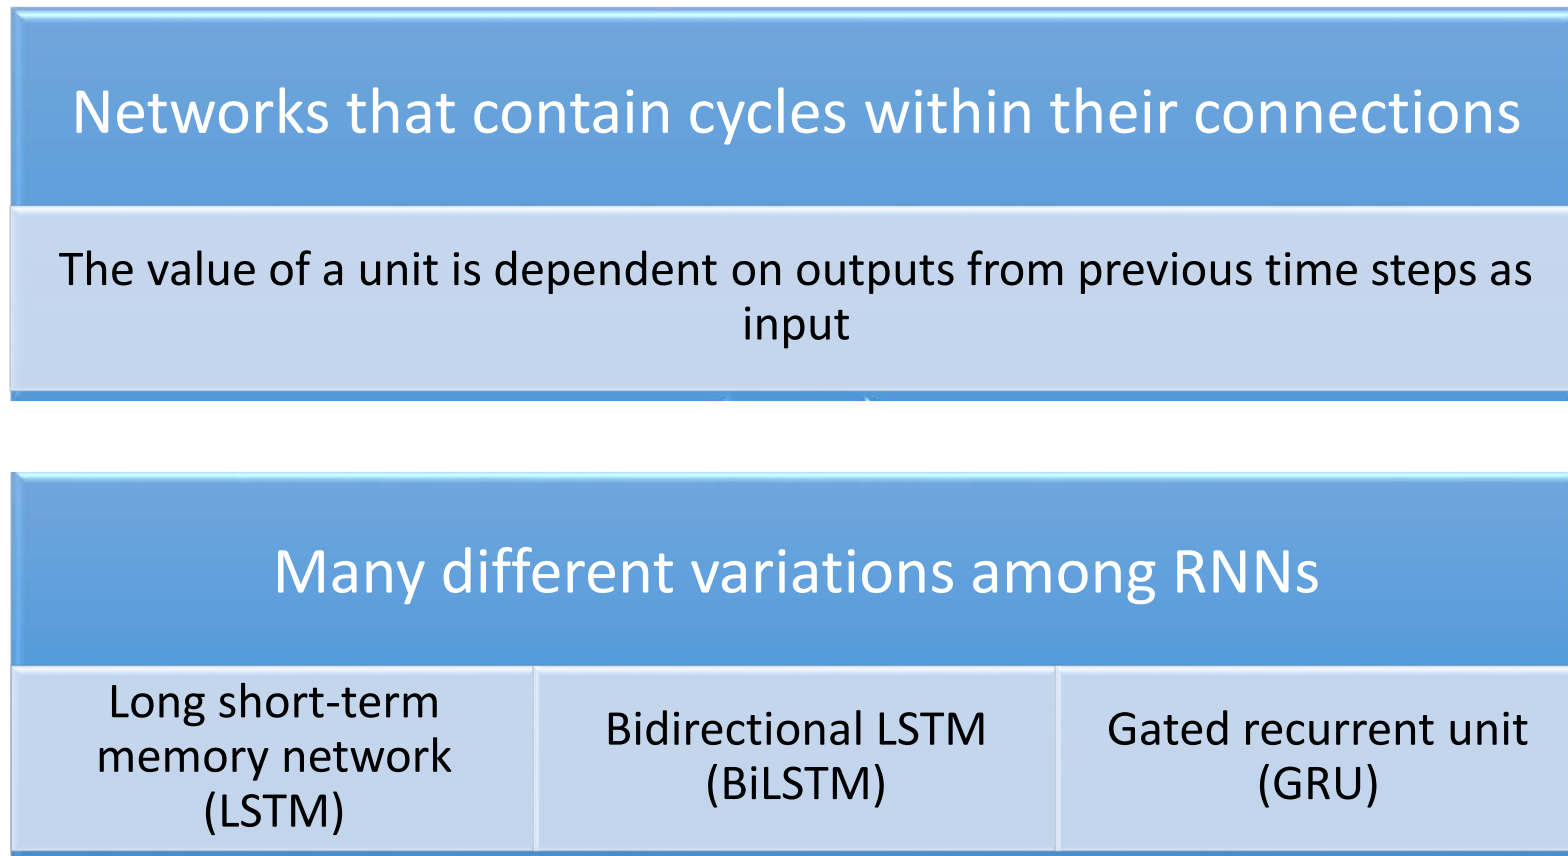
$$L = -\log P(w_t | w_{t-1})$$

- Limits the context from which
- information can be extracted
- Makes it difficult to learn systematic patterns

Sequence Processing with Recurrent Networks


Address limitation of sliding window approach

Handle variable length input



Decoding for sequence labeling

- HMM
- Linear-Chain Conditional Random Fields:
 - Viterbi algorithm for inference
- Neural model (Recurrent NN) (Elman, 1990) –
- What are the challenge?
 - input and output of the network does not have a fixed length
 - **Solution:** activation value of the hidden layer h_t depends on the current input x_t as well as the activation value of the hidden layer from the previous time step h_{t-1} .



Simple recurrent neural network illustrated as a feed-forward network




Figure 9.3 Simple recurrent neural network illustrated as a feed-forward network.

RNN unrolled in time + Inference

```
function FORWARDRNN( $x$ ,  $network$ ) returns output sequence  $y$ 
```

```
 $h_0 \leftarrow 0$ 
```

```
for  $i \leftarrow 1$  to LENGTH( $x$ ) do
```

```
   $h_i \leftarrow g(U h_{i-1} + W x_i)$ 
```

```
   $y_i \leftarrow f(V h_i)$ 
```

```
return  $y$ 
```

Figure 9.5 Forward inference in a simple recurrent network.




Figure 9.4 A simple recurrent neural network shown unrolled in time. Network layers are copied for each timestep, while the weights U , V and W are shared in common across all timesteps.

Training an RNN

- Loss Function
- Backpropagation *through time*

Two-pass algorithm for training RNNs

First pass: **Perform forward inference**

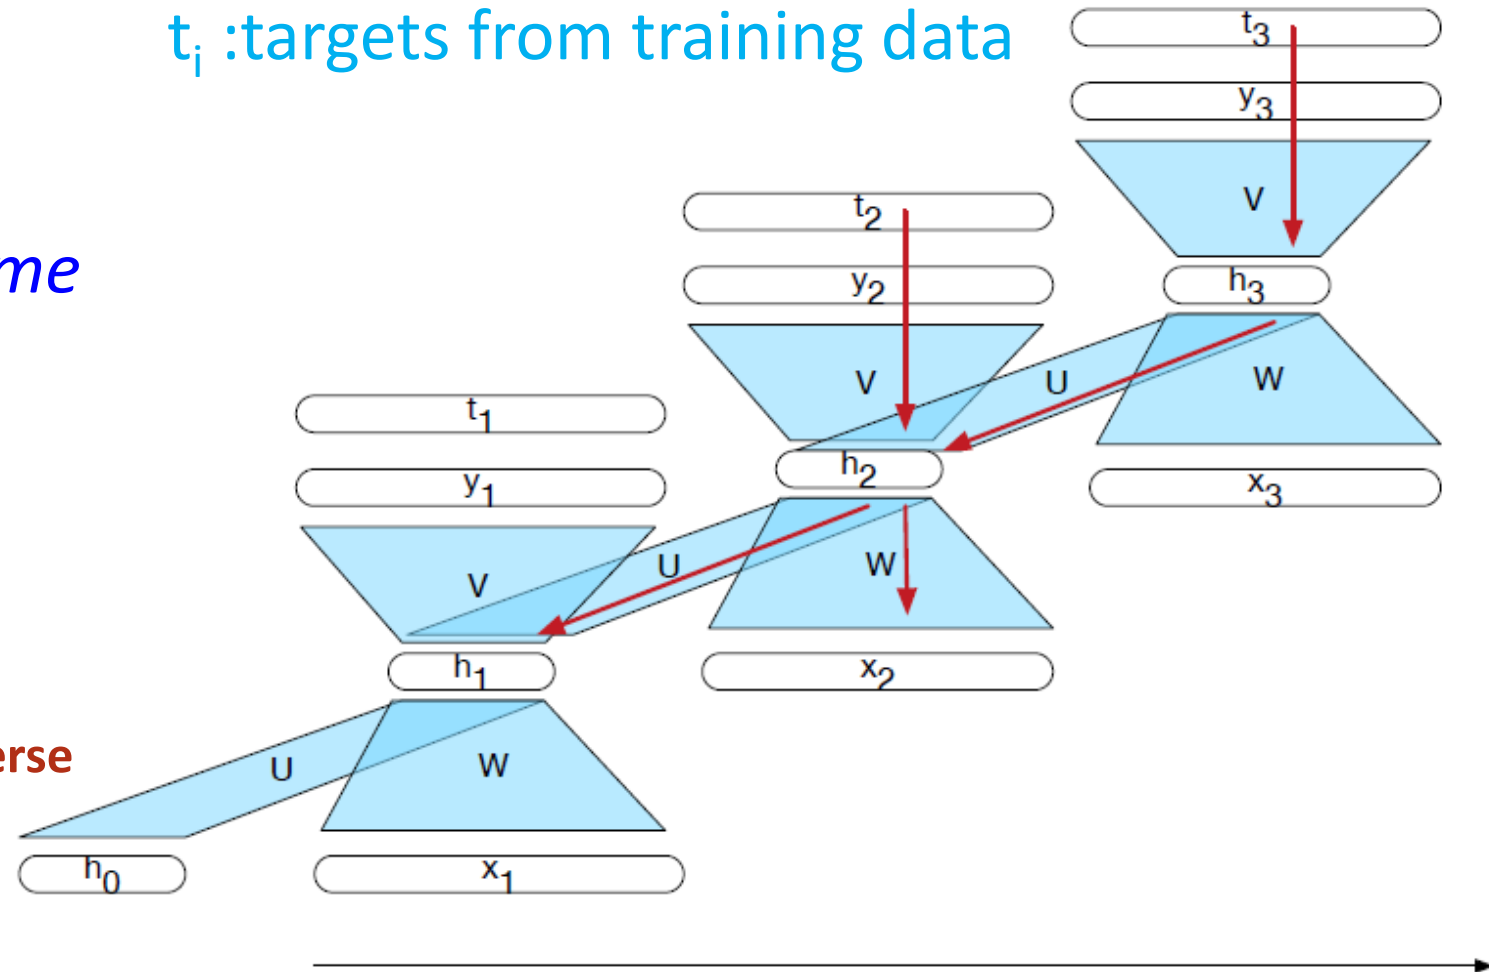
Compute h_t and y_t at each step in time

Compute the **loss** at each step in time

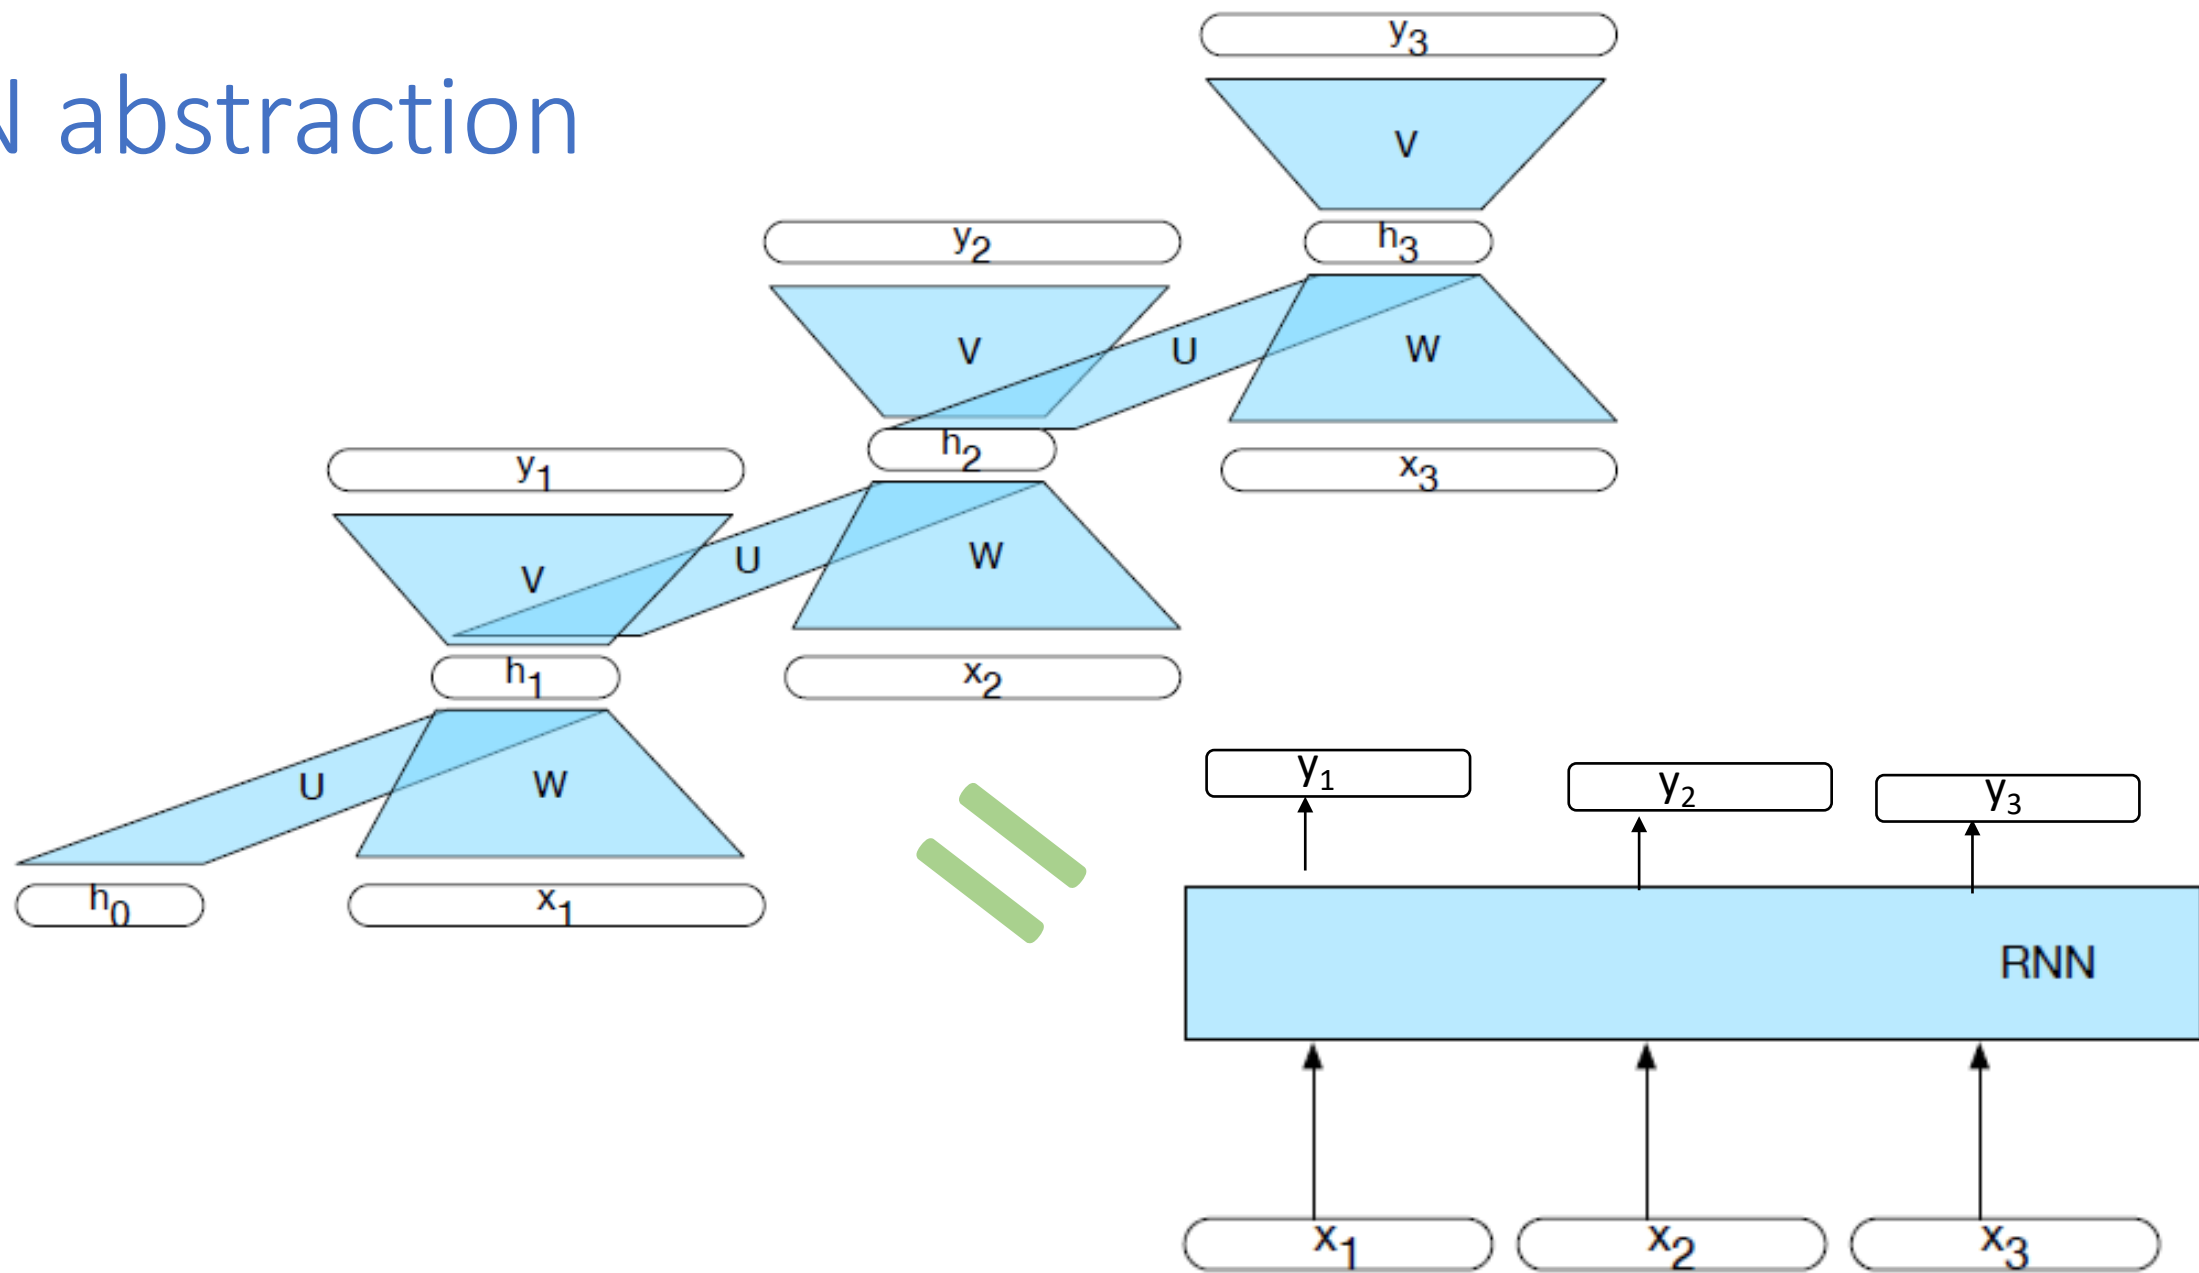
Second pass: **Process the sequence in reverse**

Compute the required error gradients at each step backward in time

t_i : targets from training data



RNN abstraction

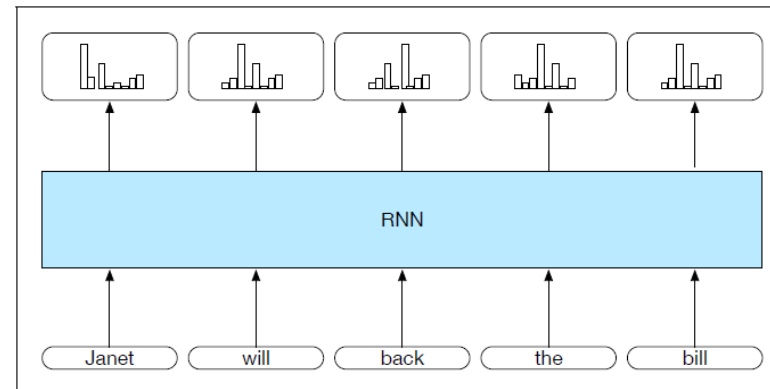
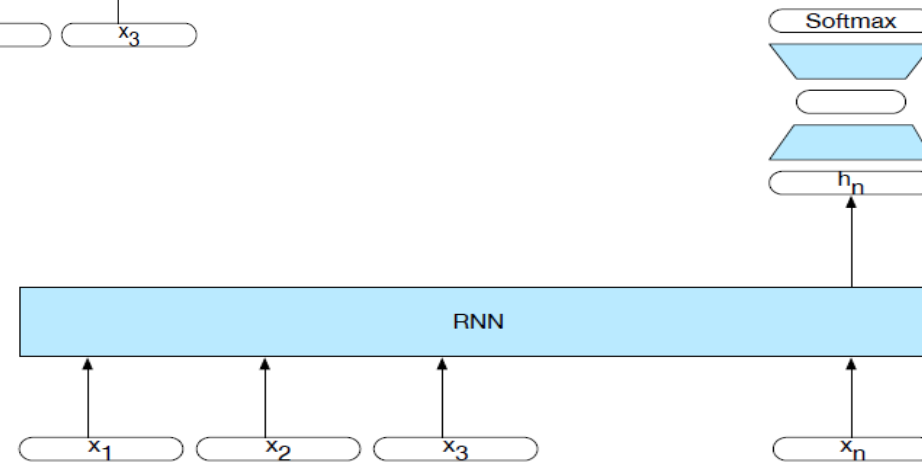
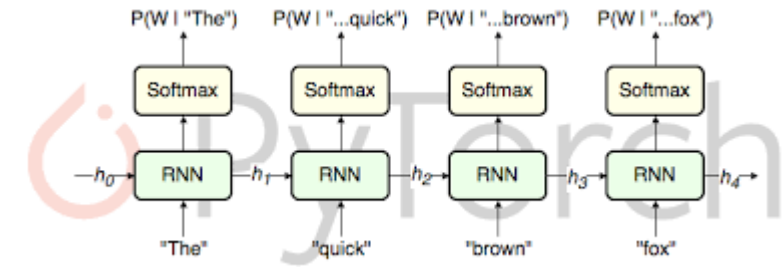
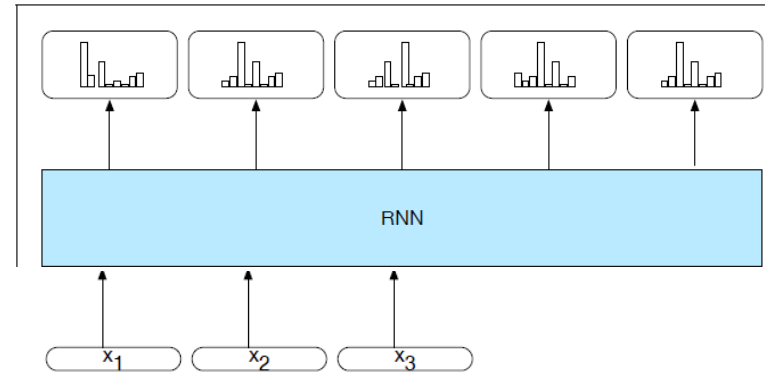


Truncated Backpropagation training through time

- For applications that involve much longer input sequences, such as speech recognition, character-by-character sentence processing, or streaming of continuous inputs, unrolling an entire input sequence may not be feasible.
- In these cases, we can unroll the input into manageable fixed-length segments and treat each segment as a distinct training item. This approach is called **Truncated Backpropagation Through Time (TBTT)**.

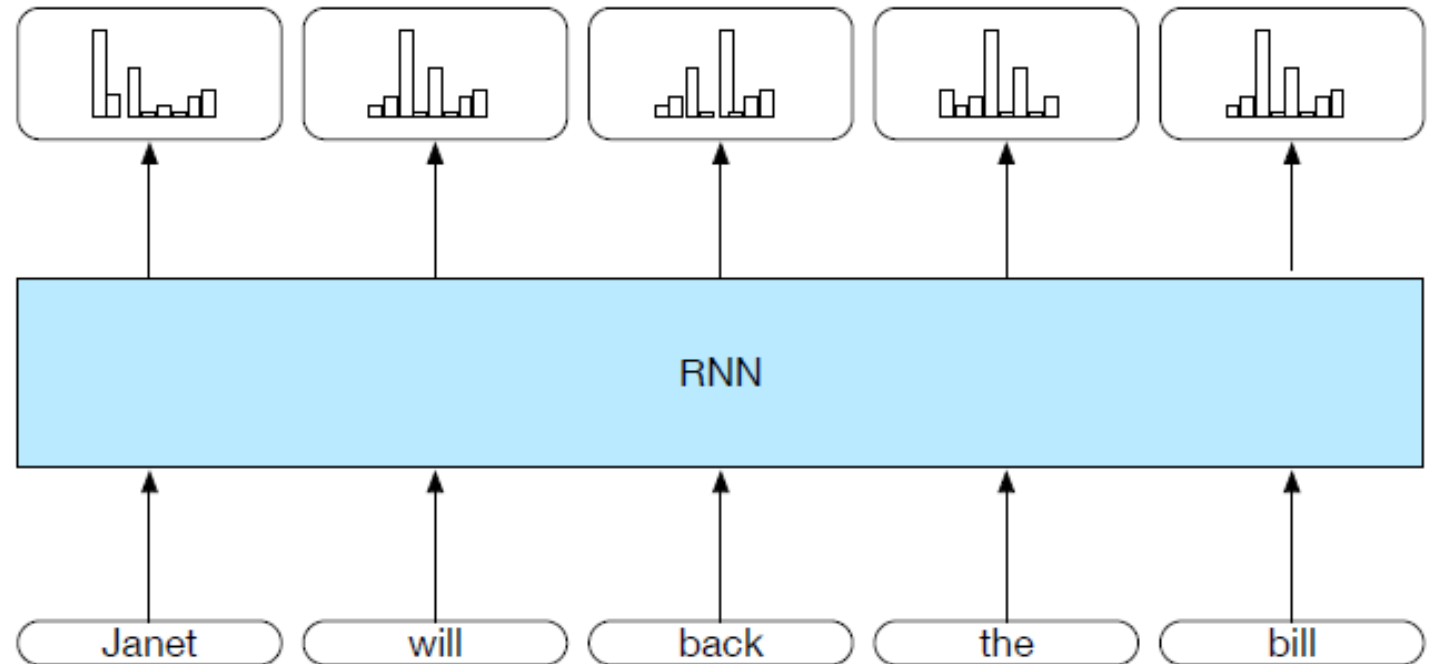
RNN Applications

- Language Modeling
- Sequence Classification (Sentiment, Topic)
- Sequence to Sequence



RNN Applications: Sequence Labeling (e.g., POS)

- **Input:** pre-trained embeddings
- **Output:** softmax layer provides a probability distribution over the part-of-speech tags as output at each time step
- Choosing max probability label for each item does not necessarily result in optimal (or even very good) tag sequence
- Combine with Viterbi for most likely sequence



RNN Applications: sequence classification (RNN + FeedForward)

- Hidden layer from final state (compressed representation of entire sequence) ->
- Input to feed-forward trained to select correct class
- No intermediate outputs for items in the sequence preceding x_n
=> no intermediate losses
- Only cross-entropy loss on final classification backpropagated all the way...

